

SparkMyName — Session Continuity & Handover

Session of 29–30 July 2026 · Founder: Peter Klein · Written for whoever picks this up next. Everything below is verified against the code, not remembered. The running ledger inside every disc (**WHAT_WAS_DONE_2026-07-29.txt**) holds the long-form detail for each item.

0 START HERE

Baseline disc	SPARK (LawnEssence standard - brand card matches the reference) July 30 2026 - 0810.zip — the last disc of this session. It contains every change listed here. Load it, unzip to <code>/home/claude/site</code> , work from there.
The ledger	WHAT_WAS_DONE_2026-07-29.txt at the root of that disc. ZIP 9 through ZIP 24. Read it before touching anything — each entry says what changed, why, and how it was proven.
Working method the Founder enforces	“Do not assume anything is true or works until you PROVE, VERIFY, TEST, and sweep for bugs — THEN issue a zip.” In practice: identify images BY EYE from disk (never trust a filename); node --check every touched JS file; drive real behaviour through jsdom rather than reading code and hoping; diff every edit against the prior file; append a ledger entry per disc; ship small and let him test between.
Deploy	The Founder deploys himself at app.netlify.com/projects/sparkmyname/deloys (NOT <code>/drop</code>). Netlify login is Gmail only.
Zip naming	NO COMMAS anywhere in a delivered zip filename , including the date — commas break his handling of the file. Write “July 30 2026”, never “July 30, 2026”.
How he works	Voice dictation — read for intent, not literal wording. ADA: one small step per message, name the exact window, button and location, then wait. He cannot right-click (use Control+Click). Hypoluxo, FL.

1 WHAT SHIPPED THIS SESSION

Disc	What it did	Proof
9	Homepage “What you get”: 8-card grid replaced with the 6-pillar grid. Three copy claims corrected to match what the engine actually returns.	—
10	Journey tabs finished: Dream / Brand / Launch / Grow, the KICK array synced, step-04 placeholder replaced with real copy.	37 assertions
11	Brand film player built. The old play button raised a toast claiming films played; there were zero video elements on the page. Films now actually play.	46 assertions
12	film-check.html — a private diagnostic page that walks the film chain and says in plain words which link is broken. No console needed.	17 assertions
13	WebP conversion: homepage images 4.75 MB → 3.16 MB. Social preview and email art deliberately left as JPEG.	12 assertions
14–17	Realtime voice intake — “Talk it through” on the homepage. Built, then twice corrected (see section 3).	24 + 9 + 15 + 23
18	The shared image library deleted. Every order now generates its own header. Both cinematic prompts given matching photographic direction.	20 assertions
19–24	Brand card rebuilt — order, duplication, layout, green, type, single column, agency presentation, then matched to the approved LawnEssence card. See section 2.	23+17+24+24+36+30

2 THE BRAND CARD — WHAT IT LOOKS LIKE NOW AND WHY

The Founder supplied two references he had approved: **SparkMyName_Card_from_real_data_CharterEdge.html** and the July-9 Vine & Crust PDF. Both are a single ~760px column, sections numbered 01–13, nothing printed over any image. The live card had drifted a long way from that. It was rebuilt to match.

Change	Why — and what it replaced
Single column	Four dated comments had accumulated on one CSS rule arguing in circles: column-count sliced sections in half → grid; align-items:start left holes → stretch; stretch inflated small tiles → back to start. All were symptoms of one cause — sections of different heights side by side. One column removes the cause and closes three complaints permanently (empty gaps, Launch posts running full width alone, cramped text). 820px, centred.
Section order	Restored to the reference: 01 Why, 02 Logo, 03 Colours, 04 Typography, 05 Voice, 06 Taglines ... About/profile drops to 10. It had drifted to About FIRST, burying the logo and palettes. One array: MERGED in js/workspace-core.js.
Duplicate removed	A real bug. The About/profile panel printed the “Why it works” bullets a second time, beside the dedicated Why tile showing the identical list. Section 01 owns them now.
Nothing over the photo	Third pass on this spot. 07-27 the scrim was deleted (“these cards carry no caption” — they did, so white text vanished). 07-30 morning I restored the scrim — wrong. The reference settled it: the name, domain and tagline already appear beneath the photo , so the overlay was a second copy laid over the thing the customer paid for. Removed. History written into the CSS.
Green eliminated	78 instances, three colours, two disguised: #F4F6F5 (x16, reads grey, is green-grey), rgba(20,60,40) (x48, tinted every hairline and muted word), rgba(18,122,64) (x14, button glow). Searching for “green” finds nothing — they were found by parsing every colour and testing hue (G > R and G > B).
Type floor	Scale started at 11px. Now starts at 12.5px; running text at 16px (the site's own body size); headings Georgia 21px with leading-zero numbers driven by a CSS counter so numbering can never drift from the order. Contrast verified: 15 colour pairings tested against WCAG, worst 6.7:1, pass mark 4.5:1.
Agency presentation	The card now wears the brand's own colours. bkAccentVars() in workspace-core reads that brand's first palette and picks a swatch that is dark enough to read (luminance < .62) and actually coloured (saturation > .12), falling back to ink if none qualifies. It sets --bk-accent / --bk-accent-soft on the card; section numbers, heading rules, the left edge of every section, bullets and group rules all take it. No two customers see the same card.
Action tiles	Groups renamed Your files / Keep building / Help & housekeeping . One primary tile in the brand colour; the destructive “Remove” stepped back to a dashed outline. All ten tiles kept and verified still wired — nothing lost.

3 REALTIME VOICE — BUILT, THEN CORRECTED TWICE

“Talk it through” sits beside the existing free browser-dictation Speak button on the homepage. The customer holds a real conversation; Spark writes the finished sentence into the existing #seed box, so the rest of the site works unchanged. Files: **netlify/functions/realtime-token.js**, **js/spark-voice.js**, **js/spark-voice-ui.js**, **css/spark-voice.css**. No new API key — it uses VITE_OPENAI_API_KEY.

Mistake	What was actually true
Session restarted every ~100s to control cost	Wrong on both counts. OpenAI already truncates long conversations server-side and the session CONTINUES. Restarting handed the customer to a model with no memory, so it re-greeted them and re-asked what they had just said — every symptom the Founder reported. It was also MORE expensive: restarting destroys the prompt cache, and cached audio is \$0.40/M against \$32/M. Fix: one session, with truncation retention_ratio 0.8 and an 8000-token ceiling — OpenAI's own documented cost control.
CSP not checked before building	netlify.toml connect-src listed Supabase and Stripe but not api.openai.com , so the browser blocked the call and reported only “Load failed”. Added, with a comment above it so a future sweep does not strip it. Lesson recorded: check connect-src FIRST whenever a feature calls a new outside service.
Input transcription never enabled	It is opt-in and billed separately, which is why the customer's own words never appeared on screen. Now on (gpt-4o-transcribe).
A “Done talking” button sat there mid-conversation	The Founder: people will press it and get cut off. Removed. Spark now closes the conversation itself after a warm goodbye. Closing early salvages whatever was understood into the idea box.

Cost, measured: gpt-realtime-2 is \$32/M audio in, \$64/M out. A five-minute intake lands around **\$0.30–\$1.00** against a \$99 sale — under 1%. No price rise is needed for this. A 20-minute wall-clock stop exists so a tab left open overnight cannot bill forever.

4 THE FILMS — ELEVEN FINISHED, THE PLAYER WAS MISSING

The Founder reported films had stopped working. **They had not.** Fifteen film manifest functions existed on the server and **zero pages called any of them** — the player had been lost in an earlier homepage rebuild, and the film wall was left as still images with a play button that raised a toast. Even the July-27 rollback homepage has no player, so it predates that point. Storage was intact throughout.

Fact	Detail
A film is not one file	Several scene .mp4s + ONE continuous narration .mp3 + ONE music .mp3. The player stitches them at playback on two alternating video elements so the next scene is already buffered.
Source of truth	The manifest, never storage. The producer prunes it to the current brand list, so orphan files (odyssey1–4, luxheli, marquee) are unreachable by construction — no filename guessing anywhere.
Eleven films exist	sparkpreview(4) jetforce(3) allies(3) romanesco(3) tresses(3) outlierdialogues(3) aurarevival(3) brightpath(3) pawsandhearts(3) odyssey5(4) sparkstudios(5). Only five are on the wall.
Captions are real	js/film-captions.js carries the exact TTS script for all 11 films, so a caption shows the words the narrator speaks. Cue TIMING is proportional by character count — wording exact, boundaries approximate. Stated in the file header.
TTS ear-check: CLOSED	Finished films cannot drift — their narration is a saved mp3. Only a producer re-run would re-record. Do not fire producer keys as a “fix”.

5 IMAGES — THE LIBRARY IS GONE, AND WHY THAT MATTERED

The Founder reported headers coming out “vintage” and mismatched — a tyre shop's picture on a dentist. Diffing the July-26 disc against live proved the cinematic prompt was **byte-for-byte identical** and had been right since 23 July. **The prompt was never the fault.** The fault was that one photograph was banked per INDUSTRY and reused forever, so a bad frame was inherited by every customer in that industry and no prompt fix could ever reach them — a library hit never regenerates. It was patched twice (restore prompt, then orphan the shelf to v2) and the pictures were still wrong.

Now: the library is deleted, not versioned again. Every order generates its own header, written to **orders/<orderkey>/header.png** — under the order's own key, never an industry name — so it can never be served to anyone else. A bad frame can only affect the single order it was made for. Cost: one generation per order instead of one per industry forever; the Founder judged that worth it, and he is right. Old library/v1 and v2 frames still sit orphaned in Supabase; nothing reads them; delete by hand if the space is wanted.

The second prompt. There genuinely were two cinematic calls. generate-asset.js carried real photographic direction the header prompt lacked — full-frame camera, fast prime lens, volumetric light, honest reflections, micro-texture. That is why the photo crowning each workspace looked flatter than the assets beneath it. Both now carry the same optics, plus a new guard: “Modern and current — photographed today, not a period or archival image.”

6 STILL OPEN

Item	Detail	Owner
Downloads need wiring properly	Founder's words: “Download this brand means EVERYTHING about this thing.” And: everything in the ZIP should be a PDF except the deliverables , and the brand sheet should print everything on the brand card, in colour . All ten tiles are wired and functional today, but the CONTENTS have not been audited against that standard. He asked only that nothing be lost for now — he will organise it later. This is the biggest open build.	Next session
film-check never run	The diagnostic page was built and shipped but the Founder never reported which step goes red. Deploy and open /film-check.html . It answers in one look whether films play.	Founder

Item	Detail	Owner
Six films invisible	Eleven finished, five on the wall. Four (allies, brightpath, outlierdialogues, pawsandhearts) already have stills in img/ — small disc. sparkpreview and sparkstudios need a still frame. Overtime Social (9 scenes, own manifest endpoint, player already supports source:'overtime2') needs one frame for a tile.	Next session
Homepage has no Open Graph tags	29 other pages have them; index.html has none. Pasting sparkmyname.netlify.app into iMessage, Slack or Facebook produces a bare link with no image or title card. Pre-existing. Small fix, real reach.	Next session
gpt-image-2 env flip	DONE — OPENAI_IMAGE_MODEL=gpt-image-2 was added in Netlify this session. gpt-image-1 shuts down 23 Oct 2026. Verified live, same endpoint, and the one gpt-image-2 limitation (no transparent backgrounds) does not apply to this request shape.	Closed
gpt-image-2 for text-bearing art	Not started. The ladder leads with Gemini for photographs. gpt-image-2's headline strength is accurate TEXT rendering — which matters for cards, signage and packaging carrying the customer's name. Proposal: photos stay on Gemini, text-bearing assets lead with gpt-image-2. Run a side-by-side first.	Next session
Web search in Market Pulse	Not started. One of the eight AI tools; currently answers from model memory so it cannot speak to a customer's market this month. Contained change.	Next session
USPTO status check	Serial No. 50014040, filed 26 July 2026, owner VORREX IGNITE LLC. Check TSDR ~late October to avoid abandonment. Use ™ only — never ® until a certificate actually issues.	Founder

7 THE TWO RULES THAT SURVIVE

The Founder revoked all standing design, layout, process and folder rules on 28 July 2026. Do not cite retired rules back at him and do not apply them silently — he dictates as he goes. **Two things are not style preferences and stay:**

- 1. Truth** Never state as fact something the software does not do. This protects him from refunds and complaints. It is why three copy claims were corrected on the homepage this session (three palettes not five, six launch posts not a month, the five real vector files not a "reversed" one), why the film player's failure message never pretends it heard anything, and why the voice persona is forbidden from promising websites, hosting, or a guaranteed domain.
- 2. No ®** Never use ® before a registration certificate actually issues — it is unlawful and can jeopardise the pending application. ™ is correct and is used everywhere the wordmark appears. Keep both out of domains, email addresses and machine-readable SEO data.

One habit worth keeping. Several mistakes this session shared a shape: building before checking. The rollover was invented before OpenAI's cost guidance was read. The voice feature was built before the site's own CSP was checked. The photo scrim was restored before the Founder's reference card was asked for. In each case the answer already existed and one look would have found it. **Read the vendor's own documentation, and ask for the reference, BEFORE designing around a problem.**

Also worth knowing: several test failures this session were the TEST's fault, not the site's — a regex matching a comment that contained the very word being searched for, an assertion assuming the homepage had an og:image when it does not. When an assertion fails, check the test before changing the code. And a claim was once written into the ledger before the suite had actually run; it was caught and corrected. **The ledger is only worth something if every line in it is true.**